

QuietBox 2 LDPC Decoder Engine: Architecture & Scaling Roadmap

Autonomous On-Device Monte Carlo Simulation down to 10^{-8} – 10^{-10} Error Floor

ALGORITHM ORIGINATOR	HARDWARE PLATFORM	PRIMARY TARGET CODE	OPERATIONAL DRIVER
Dr. Christopher R. Jones (JPL / MILCOM 2003)	QuietBox 2 (4 Blackhole ASICs, 2x2 Mesh)	CCSDS AR4JA Rate-1/2 (N=2560, M=1536)	TT-Metal 0.54 / tt-kmd 2.11.0

1. Executive Summary


This technical report details the architectural foundation and high-throughput execution model of the **QuietBox 2 (QB2) LDPC Monte Carlo Simulator** deployed across 440 Blackhole Tensix cores. The mission of this engine is to simulate billions of codewords to explore the deep-space communication error floor ($10^{-8} \leq FER \leq 10^{-10}$) completely autonomously on-device, bypassing host PCIe streaming bottlenecks.

The core decoding algorithm employs **Dr. Christopher Jones' Approximate-Min* algorithm (MILCOM 2003)** using a row-layered Gauss-Seidel schedule, native `bfloat16` precision, and inline early parity syndrome detection. This report reviews recent experimental breakthroughs—including the formal resolution of the 2.0 dB "knee" artifact and characterization of the 1.75×10^{-8} floor—and lays out a phased engineering roadmap to scale throughput from **46 Msps to over 2.0 Gsps** across the mesh.

2. Hardware Topology & Compute Hierarchy

The QuietBox 2 is a quiet, liquid-cooled workstation featuring **4 Tenstorrent Blackhole processors** interconnected in a 2x2 Torus/Mesh fabric:

- **Mesh Aggregate:** 4 chips $\times$ 110 compute Tensix cores = **440 Tensix compute cores** operating at ~1.5 GHz.
- **Per-Core L1 SRAM:** 1.5 MB high-speed local memory (660 MB aggregate across the mesh) tightly coupled to 5 embedded RISC-V processors.
- **Compute TRISC Trio:** Each Tensix core contains 3 dedicated execution processors: TRISC0 (UNPACK), TRISC1 (MATH), and TRISC2 (PACK), alongside 2 Data Movement processors (BRISC, NCRISC) and a 32-wide SIMD Special Function Processing Unit (SFPU).



3. Autonomous On-Device Pipeline & Precision Engineering

Simulating deep error floors (10^{-8} to 10^{-10} FER) requires evaluating hundreds of millions to billions of codewords. Streaming noise frames and decoding results over PCIe would limit throughput to under 5 Msps. The QB2 engine executes the entire lifecycle inside Tensix L1 SRAM:

1. **Polar Box-Muller AWGN Generator:** Evaluates standard normal variates using rejection sampling in the unit disk ($s = u^2 + v^2 < 1.0$), generating pairs of independent Gaussian random variables without trigonometric calculations (sin/cos).
2. **Zero-Overhead Puncturing:** For the AR4JA rate-1/2 code, $P = 512$ systematic parity bits are not transmitted. Their channel observations are initialized to neutral LLR ($L_i = 0$) directly in L1.
3. **Row-Layered Convergence:** Variable node beliefs L_v are updated immediately in place after each check row is processed, reducing convergence iterations by $2\times$ over two-phase flooding.
4. **Inline Parity Termination:** After each layer, hard decisions $c_v = (L_v < 0)$ are evaluated against the parity equations ($H \cdot c^T = 0$). Valid codewords terminate instantly.
5. **PRNG Spatial & Temporal Isolation (xoshiro256**):** The 440 cores are separated by 2^{96} PRNG draws using successive `long_jump()` calls ($\sim 7.92 \times 10^{28}$ draws/core). Batches within each core advance by 2^{64} calls via `jump()`, guaranteeing zero overlap.
6. **Native Bfloat16 Precision (8x SRAM Compression):** LLRs and check messages $r_{\{m,d\}}$ are stored in native `bfloat16` (`Float16_b`). Core memory footprint dropped from **448 KB to 56 KB**, saving 172.5 MB of L1 SRAM across the mesh with identical error-correction performance.

4. Mathematical Formulation: Christopher Jones Approximate-Min*

Exact belief propagation at a check node requires evaluating Jacobi logarithms across all incoming extrinsic messages:

$$\min^*(a, b) = \ln(e^{\ln(-a)} + e^{\ln(-b)}) = \min(a, b) - \ln(1 + e^{\ln(-|a| - |b|)}) = \min(a, b) - g(|a| - |b|)$$

Dr. Christopher Jones' Approximate-Min* algorithm (MILCOM 2003) proves that full belief propagation does not require evaluating pairwise Jacobi trees for every outgoing edge. Instead:

- **3 Smallest Magnitudes:** The check node tracks only $\min_1 \leq \min_2 \leq \min_3$ and computes two difference terms: $\Delta_1 = \min_2 - \min_1$ and $\Delta_2 = \min_3 - \min_2$.
- **Only Two Outgoing Magnitudes:**

$$r_{\{\text{mag}, \text{all}\}} = \max(0, \min_1 - g(\Delta_1)) \text{ for all non-minimum edges } d \neq d_{\min}$$

$$r_{\{\text{mag}, \text{min}\}} = \max(0, \min_2 - g(\Delta_2)) \text{ for the minimum edge } d = d_{\min}$$
- **2-Region Piecewise Linear Model:** Transcendental correction $g(\Delta) = \ln(1 + e^{\ln(-\Delta)})$ is approximated using hardware shift-add / FMA friendly segments without lookup tables:

$$g(\Delta) \approx 0.69315 - 0.4000 \cdot \Delta \text{ (if } \Delta < 0.5)$$

$$g(\Delta) \approx 0.49315 - 0.3280 \cdot (\Delta - 0.5) \text{ (if } 0.5 \leq \Delta < 2.0)$$

$$g(\Delta) = 0.0 \text{ (if } \Delta \geq 2.0)$$

This eliminates lookup table contention and achieves a **57.3% frame error reduction** over Normalized Min-Sum (NMS, $\alpha = 0.75$).

5. Experimental Discoveries & Deep Floor Characterization




Figure 1: Multi-Iteration Waterfall Performance of Approximate-Min* on Tenstorrent QuietBox 2 across 440 Blackhole Tensix Cores. Comparing 16 iters, 100 iters, 200 iters with fixed $L_{\max}=4.2$ clipping, and 200 iters unclipped down to 2.40 dB (915.6M blocks evaluated).

5.1 Resolution of the 2.0 dB "Knee" Artifact

Earlier simulation sweeps employing fixed channel clipping ($L_{\max} = 4.2$) exhibited an abrupt $15.4\times$ drop between 2.00 dB and 2.10 dB. Evaluating natural unclipped AWGN proved that this knee was a clipping-induced distortion:

- Under unclipped Gaussian noise, the waterfall slope is smooth and continuous: 1.9 → 2.0 dB drops by 4.6×, 2.0 → 2.1 dB drops by 4.8×, and 2.1 → 2.2 dB drops by 4.0×.
- Fixed 4.2 clipping artificially saturated correct variable beliefs at lower SNRs. Removing it restored true theoretical log-linear descent.

5.2 Ultra-Deep Error Floor at 2.40 dB (1.75 × 10⁻⁸)

At 2.40 dB, the 200-iteration engine evaluated **915,596,000 blocks** (> 1.87 Trillion bits):

Eb/N0	Blocks Run	Bit Errors	Block Errors	BER	FER	Bits/FE	Convergence Status
2.00 dB	1,000,000	3,047	12	1.49 × 10 ⁻⁶	1.20 × 10 ⁻⁵	253.9	Waterfall
2.10 dB	9,444,000	5,992	25	3.10 × 10 ⁻⁷	2.65 × 10 ⁻⁶	239.7	Waterfall
2.20 dB	37,868,000	6,346	25	8.18 × 10 ⁻⁸	6.60 × 10 ⁻⁷	253.8	Waterfall
2.30 dB	431,698,000	6,269	25	7.10 × 10 ⁻⁹	5.79 × 10 ⁻⁸	250.8	Waterfall / Floor Transition
2.40 dB	915,596,000	4,228	16	2.60 × 10 ⁻⁹	1.75 × 10 ⁻⁸	264.3	Deep Floor Onset

Critical Insight: Non-Convergence vs. Trapping Sets

The steady average of ~264 bit errors per failed frame proves these failures are **non-convergence states** rather than classical trapping sets (which flip only 8–16 bits). In 550 billion Gaussian draws, extreme 6σ–7σ tail spikes produce initial channel LLRs of L_{ch} < -12.4 that degree-6 check nodes cannot overturn. Relaxed proportional clipping (γ = L_{max} / μ_{LLR} ≈ 2.0–2.5, L_{max} ≈ 7.0–8.5) bounds extreme noise spikes without saturating the signal distribution.

6. Bottleneck & Hardware Utilization Analysis

Profiling current single-TRISC execution across the 440-core QuietBox 2 mesh reveals extraordinary untapped headroom:

Subsystem / Hardware Component	Current Operational State	Full Tensix Capability	Current Silicon Utilization
Compute TRISC Processors	Only TRISC0 (UNPACK) active	TRISC0 + TRISC1 + TRISC2 concurrent	33.3% (66.7% idle)
Data Movement (BRISC / NCRISC)	Idle during compute iterations	Autonomous background DM & PRNG	< 1%
Special Function Processing Unit (SFPU)	Unused (scalar C++ code)	32-wide SIMD vector math unit	0% (100% idle)
Per-Core L1 SRAM Memory	56 KB utilized	1,500 KB available	3.7% (96.3% headroom)
Mesh Simulation Throughput	46.2 Msps (16-iter baseline)	Scalable to > 2,000 Msps	Baseline

7. Throughput Scaling Roadmap

To evaluate billions of codewords at 2.4–2.6 dB in hours rather than days, we establish a four-phase acceleration roadmap:

[Current Baseline: 46 Msps]

▼ (Phase 1: Multi-TRISC Concurrency)

[138 Msps (3.0x Speedup) across 440 Cores]

▼ (Phase 2: L1 Double Buffering)

[170 Msps (3.7x Speedup)]

▼ (Phase 3: SWAR Register-Level SIMD)

[550 Msps (12.0x Speedup)]

▼ (Phase 4: 32-Wide Tensix SFPU Vectorization)

[1.8 - 2.5 Gsps (40x - 55x Speedup)]

Phase 1: Multi-TRISC Concurrency (TRISC0 + TRISC1 + TRISC2)

- **Target Speedup: 3.0×** (~46 Msps → **138 Msps**).
- **Implementation:** Partition L1 SRAM into three non-overlapping 56 KB blocks (3 × 56 KB = 168 KB, taking only 11.2% of L1). Launch TRISC0, TRISC1, and TRISC2 concurrently on independent codewords with disjoint PRNG seeds.
- **Effort:** Low–Medium (1–2 days). Zero hardware risk.

Phase 2: L1 Ping-Pong Double Buffering

- **Target Speedup: +20% to 25%** (cumulative ~**170 Msps**).
- **Implementation:** Decouple Box-Muller AWGN generation from LDPC iteration loops using ping-pong buffers (Buffer A actively decoded while Buffer B generates next-batch noise). Eliminates the 18–22% PRNG generation stall.
- **Footprint:** 336 KB (22.4% of L1 SRAM).

Phase 3: SWAR (SIMD Within A Register) Parallelism

- **Target Speedup: 3.0× to 4.0×** (cumulative ~**550 Msps**).
- **Implementation:** Pack two or four `bfloat16` messages into 32-bit/64-bit integer registers. Replace branchy minimum-finding with bitwise parallel min/max operations across multiple codewords in lockstep.

Phase 4: Tensix SFPU 32-Wide Vectorization

- **Target Speedup: 10× to 15×** over Phase 1 (**1.8 to 2.5 Gbps aggregate mesh throughput**).
- **Implementation:** Map Box-Muller generation and check node updates to the 32-lane Tensix SFPU vector instructions. At 2.0 Gbps, simulating **1 Billion codewords** (2.05×10^{12} bits) will complete in **under 15 minutes**.

8. Engineering Milestones & Verification Plan

Milestone	Key Deliverable	Mesh Throughput	1B Blocks Time	Status
M1	Single-TRISC BF16 Approximate-Min* on 440 Cores	46 Msps	10.5 Hours	Complete & Verified
M2	Multi-TRISC Concurrency (TRISC0 + TRISC1 + TRISC2)	138 Msps	3.5 Hours	Ready for Deployment
M3	Proportional Gamma Clipping Sweep ($\gamma = 2.0\text{--}2.5$)	138 Msps	3.5 Hours	Scheduled next
M4	L1 Ping-Pong Double Buffering	170 Msps	2.8 Hours	Planned
M5	SWAR Bfloat16 Register Parallelism	550 Msps	53 Minutes	Planned
M6	32-Wide Tensix SFPU Vectorization (Deep Floor 10^{-10})	2,000+ Msps	< 15 Minutes	Planned

References: C. Jones et al., "Approximate-Min* constraint node updating for LDPC and turbo decoding," IEEE MILCOM 2003. • CCSDS 131.0-B-5 TM Synchronization and Channel Coding (AR4JA Rate-1/2). • Tenstorrent TT-Metalium SDK 0.54 & tt-kmd 2.11.0.